

Exorcising Spectres with Secure Compilers – CSC Report



Marco Patrignani^{1,2}

23rd June 2021



CISPA
HELMHOLTZ CENTER FOR
INFORMATION SECURITY

²



Stanford
University

Talk Outline

Who Am I ?

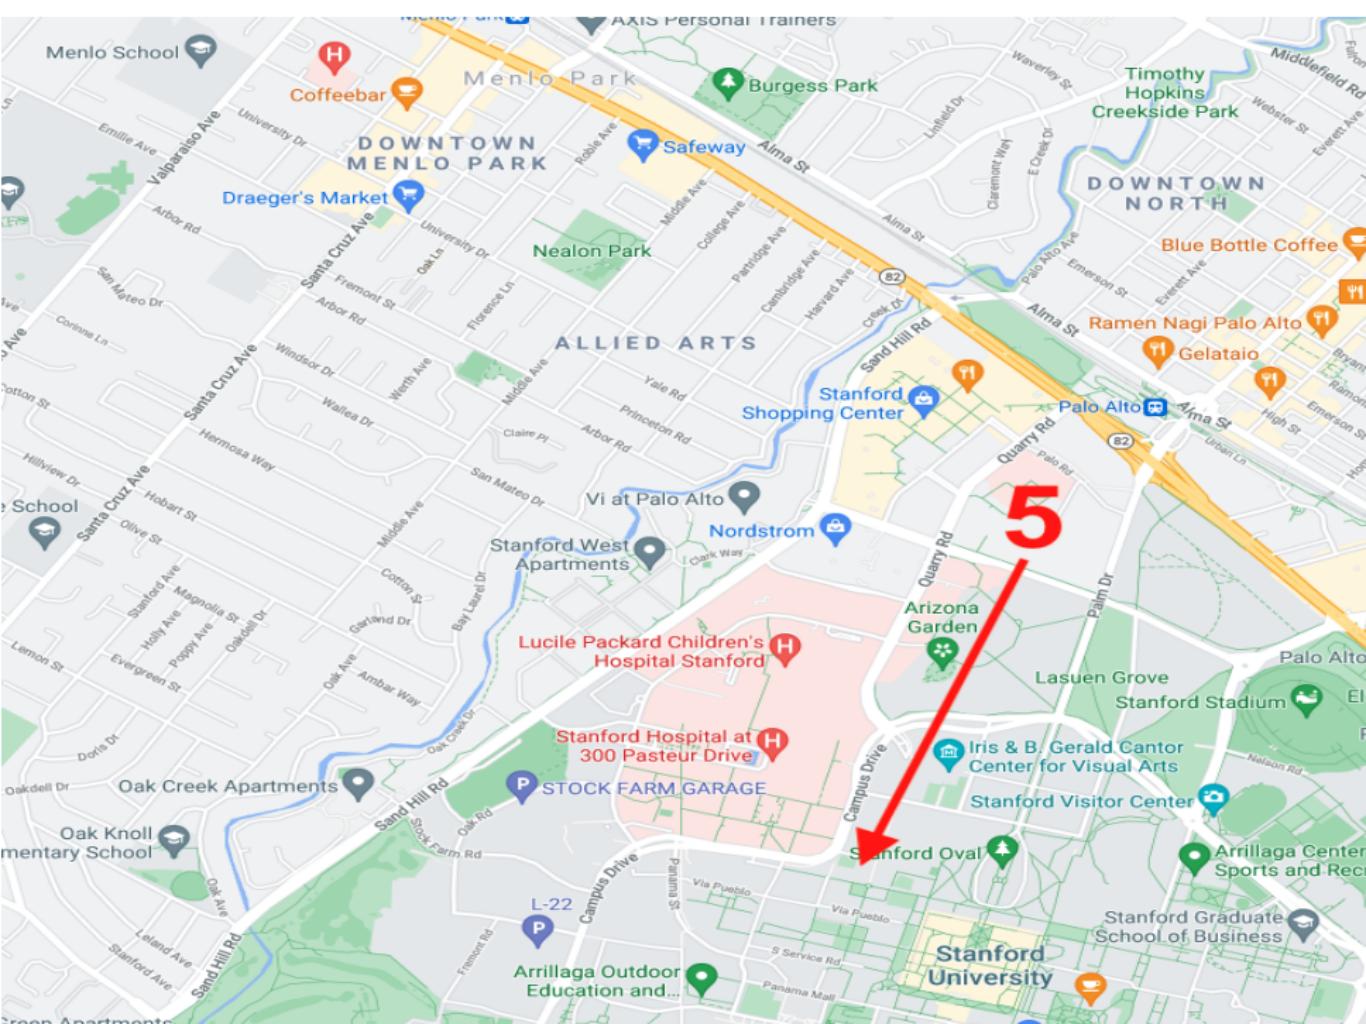
Exorcising Spectres with Secure Compilers Patrignani &

Guarnieri CCS'21

Future Outlook

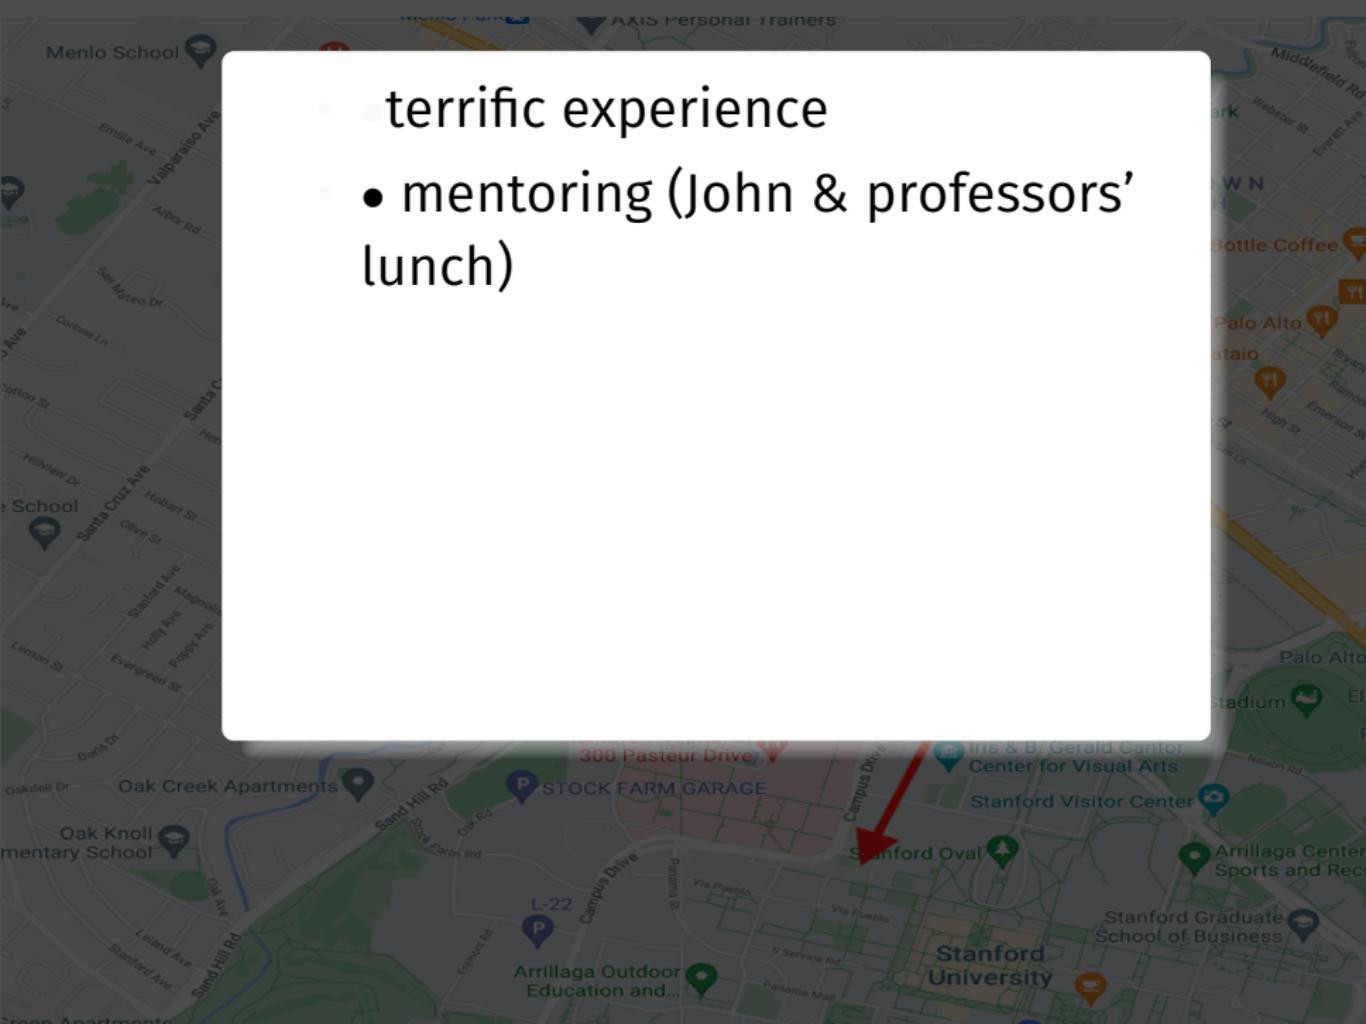
Who Am I ?





terrific experience

- mentoring (John & professors' lunch)



terrific experience

- mentoring (John & professors' lunch)
- teaching experience (courses & lunches)

terrific experience

- mentoring (John & professors' lunch)
- teaching experience (courses & lunches)
- research (+ talks)

terrific experience

- mentoring (John & professors' lunch)
- teaching experience (courses & lunches)
- research (+ talks)
- new perspective

terrific experience

- mentoring (John & professors' lunch)
- teaching experience (courses & lunches)
- research (+ talks)
- new perspective
- skiing (who'd have thought?)

Exorcising Spectres with Secure Compilers

Patrignani & Guarnieri CCS'21

Speculative execution + branch prediction

Size of array **A**

```
if (x < A_size)
    y = B[A[x]]
```

Speculative execution + branch prediction

Size of array **A**

```
if (x < A_size)
    y = B[A[x]]
```

Speculative execution + branch prediction

Size of array **A**

```
if (x < A_size)   
y = B[A[x]]
```



Branch predictor

Speculative execution + branch prediction

```
if (x < A_size) {  
    y = B[A[x]]  
}
```

Size of array **A**

Prediction based on **branch history** & **program structure**



Branch predictor

Speculative execution + branch prediction

```
if (x < A_size)  
    y = B[A[x]]
```

Size of array **A**

Prediction based on **branch history** & **program structure**



Branch predictor

Speculative execution + branch prediction

```
Size of array A  
if (x < A_size)  
    y = B[A[x]]
```

Prediction based on **branch history** & **program structure**



Wrong prediction? **Rollback changes!**



Architectural (ISA) state



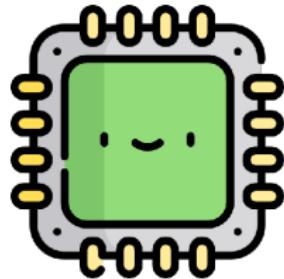
Microarchitectural state

Spectre V1

```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```

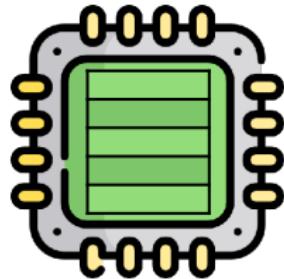
Spectre V1

```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



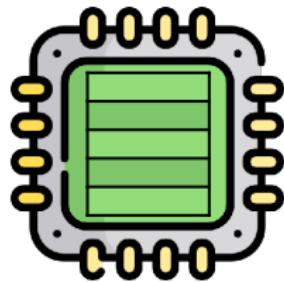
Spectre V1

```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



Spectre V1

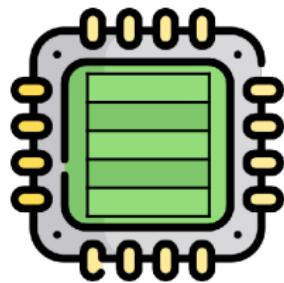
```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



Spectre V1



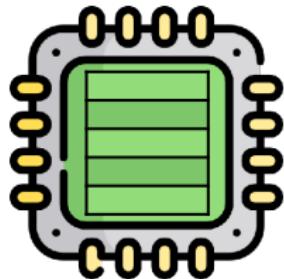
```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



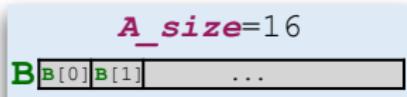
Spectre V1



void *f*(int *x*)
if (*x* < *A_size*)
y = *B[A[x]]*



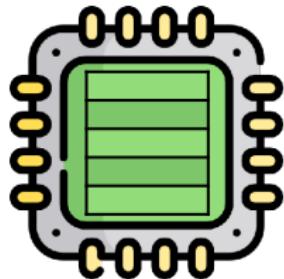
Spectre V1



```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



What is in A[128]?



Spectre V1

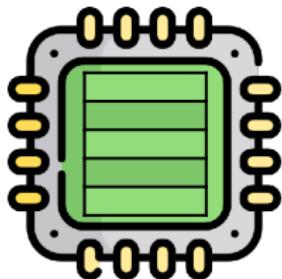


What is in **A**[128]?



```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```

1) Training



Spectre V1

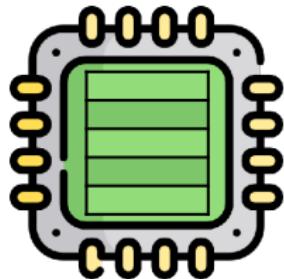


What is in **A**[128]?

```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



1) Training



Spectre V1

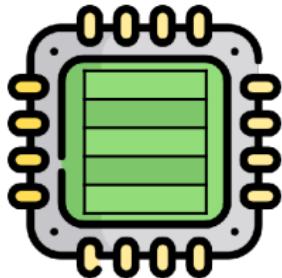


```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



What is in A[128]?

1) Training ↗ f(0);



Spectre V1



```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```

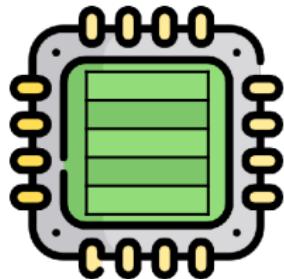
A_size=16

B[B[0] B[1] ...]



What is in **A**[128]?

1) Training f(0);f(1);



Spectre V1

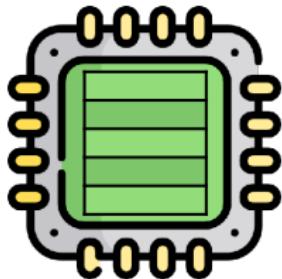


```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



What is in A[128]?

1) Training f(0);f(1);f(2); ...



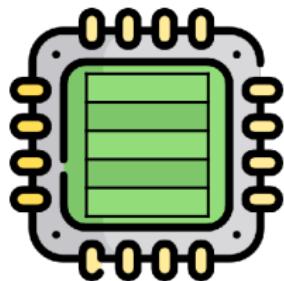
Spectre V1



```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



What is in **A**[128]?



1) Training f(0);f(1);f(2); ...

2) Prepare cache

Spectre V1



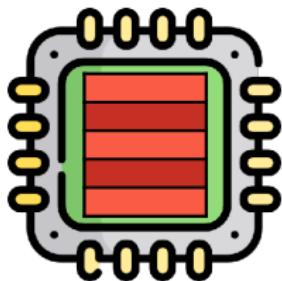
What is in **A**[128]?

```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



1) Training f(0); f(1); f(2); ...

2) Prepare cache



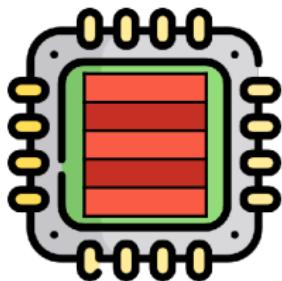
Spectre V1



```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



What is in **A[128]**?



1) Training $f(0); f(1); f(2); \dots$

2) Prepare cache

3) Run with $x = 128$

Spectre V1

The diagram illustrates the Spectre V1 attack flow through three main stages:

- 1) Training**: Shows the function `f(int x)` with the condition `(x < A_size)`. A call to `f(0); f(1); f(2); ...` is shown. A haloed smiley face icon is associated with this stage.
- 2) Prepare cache**: Shows the variable `A_size = 16` and memory layout for arrays `B` and `A`. It asks "What is in `A[128]`?". A blue ghost icon is associated with this stage.
- 3) Run with `x = 128`**: Shows the final step of running the function with `x = 128`.

Below the stages is a stylized illustration of a computer chip with red internal components and yellow pins.

7

Spectre V1



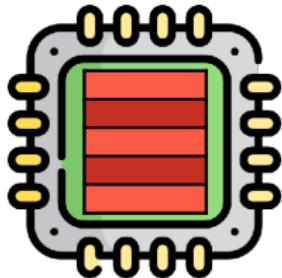
```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```

A_size=16

B[B[0] B[1] ...]



What is in *A*[128]?

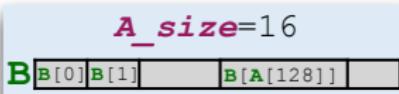


1) Training f(0); f(1); f(2); ...

2) Prepare cache

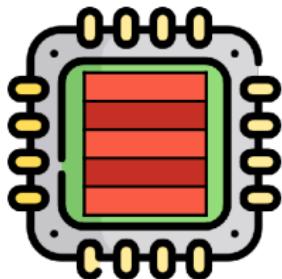
3) Run with *x* = 128

Spectre V1



What is in **A[128]**?

```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



1) Training f(0); f(1); f(2); ...

2) Prepare cache

3) Run with **x** = 128

Spectre V1



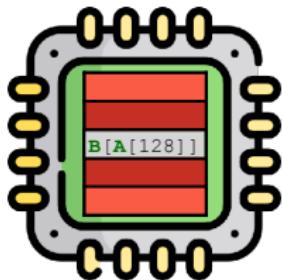
```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```

A_size=16

B[B[0] B[1] ... B[A[128]]]



What is in **A**[128]?



1) Training f(0); f(1); f(2); ...

2) Prepare cache

3) Run with *x* = 128

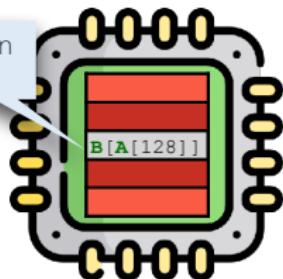
Spectre V1



```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



What is in $\mathbf{A}[128]$?

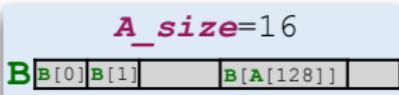


1) Training $f(0); f(1); f(2); \dots$

2) Prepare cache

3) Run with $x = 128$

Spectre V1

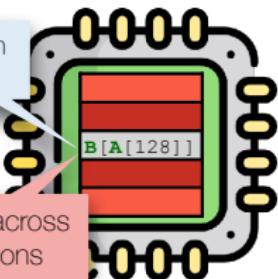


What is in $A[128]$?

```
void f(int x)
if (x < A_size)
    y = B[A[x]]
```



Depends on
 $A[128]$



Persistent across
speculations

1) Training $f(0); f(1); f(2); \dots$

2) Prepare cache

3) Run with $x = 128$

Spectre V1

The diagram illustrates the Spectre V1 attack flow through four numbered steps:

- 1) Training**: Shows the function `f(int x)` with code `if (x < A_size) y = B[A[x]]`. A haloed smiley face icon is associated with this step.
- 2) Prepare cache**: Shows a memory layout for variable `B` with `A_size = 16`. It contains elements `B[0], B[1], ..., B[A[128]]`.
- 3) Run with $x = 128$** : Shows a processor chip with a speech bubble stating "Depends on `A[128]`". A pink box below states "Persistent across speculations".
- 4) Extract from cache**: Shows a blue ghost-like character pointing to the element `B[A[128]]` in the memory layout, with the question "What is in `A[128]`?".

Compiler-level countermeasures

Compiler-level countermeasures

For *Spectre V1*

Injecting speculation barriers

```
if (x < A_size)  
    y = B[A[x]]
```



```
if (x < A_size)  
    lfence  
    y = B[A[x]]
```

- In x86, **LFENCE** act as **speculation barrier**
- Compiler injects LFENCE after each branch instruction
 - Microsoft Visual C++
 - Intel ICC
- Effectively **stop speculative execution!**

Speculative load-hardening (SLH)

```
if (x < A_size)  
    y = B[A[x]]
```



```
if (x < A_size)  
    y = B[mask(A[x])]
```

- Injects ***data dependencies*** and ***masking operations***
- Combines ***conditional moves*** and ***binary operations***
- Stops ***speculative leaks***
- Does not block speculative execution!
- Implemented in Clang

Goal

1. formalise lfence & SLH compilers

Goal

1. formalise lfence & SLH compilers
2. **T** must capture speculative execution (\rightsquigarrow)

Goal

1. formalise lfence & SLH compilers
2. **T** must capture speculative execution (\rightsquigarrow)
3. need a safety property capturing vulnerability to Spectre v1: *SS*

Goal

1. formalise lfence & SLH compilers
2. **T** must capture speculative execution (\rightsquigarrow)
3. need a safety property capturing vulnerability to Spectre v1: *SS*
4. use secure compilation to preserve *SS*: RSSC

Goal

1. formalise lfence & SLH compilers
2. **T** must capture speculative execution (\rightsquigarrow)
3. need a safety property capturing vulnerability to Spectre v1: *SS*
4. use secure compilation to preserve *SS*: RSSC
5. prove the compilers attain RSSC

Goal Up Next

2. **T** must capture speculative execution (\rightsquigarrow)
3. need a safety property capturing vulnerability to Spectre v1: *SS*
4. use secure compilation to preserve *SS*: RSSC

Speculative Semantics 101

“Spectector ...” S&P'20

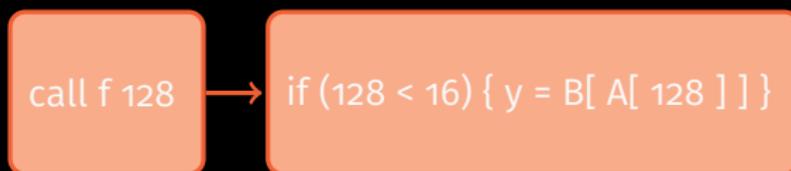
```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] }      // A.size=16, A[128]=3
```

call f 128

Speculative Semantics 101

“Spectector ...” S&P'20

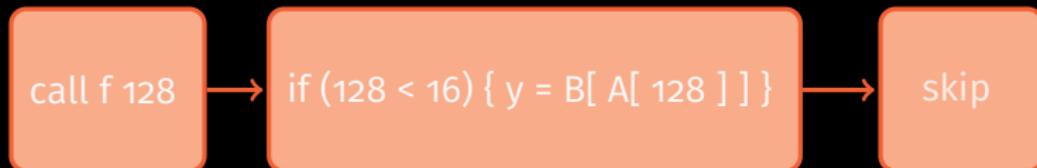
```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] }      // A.size=16, A[128]=3
```



Speculative Semantics 101

“Spectector ...” S&P'20

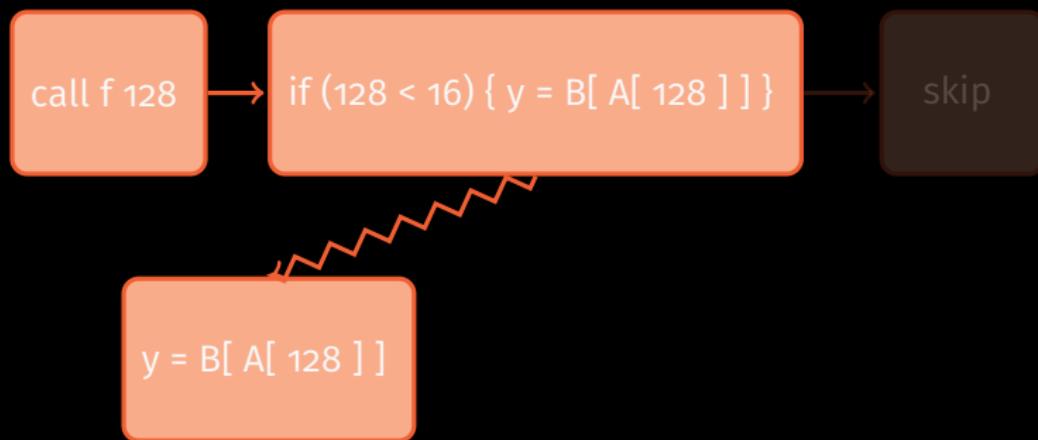
```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] }      // A.size=16, A[128]=3
```



Speculative Semantics 101

"Spectector..." S&P'20

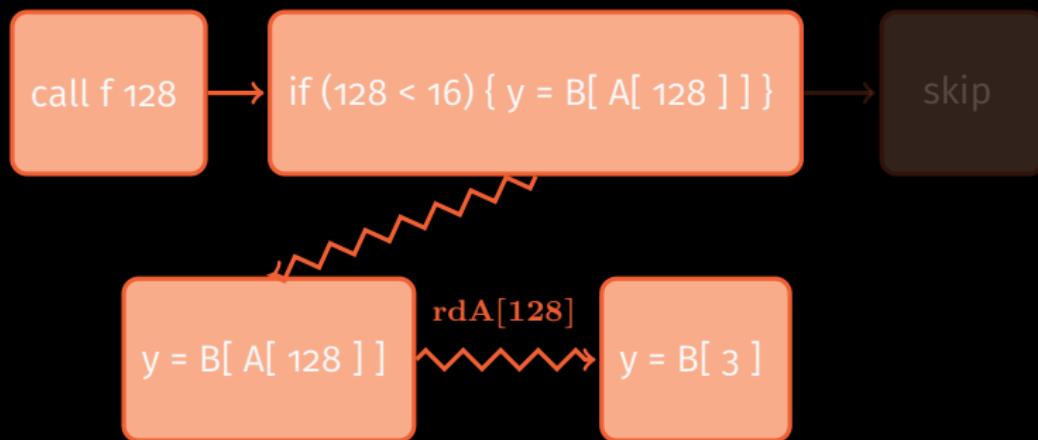
```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] }      // A.size=16, A[128]=3
```



Speculative Semantics 101

"Spectector..." S&P'20

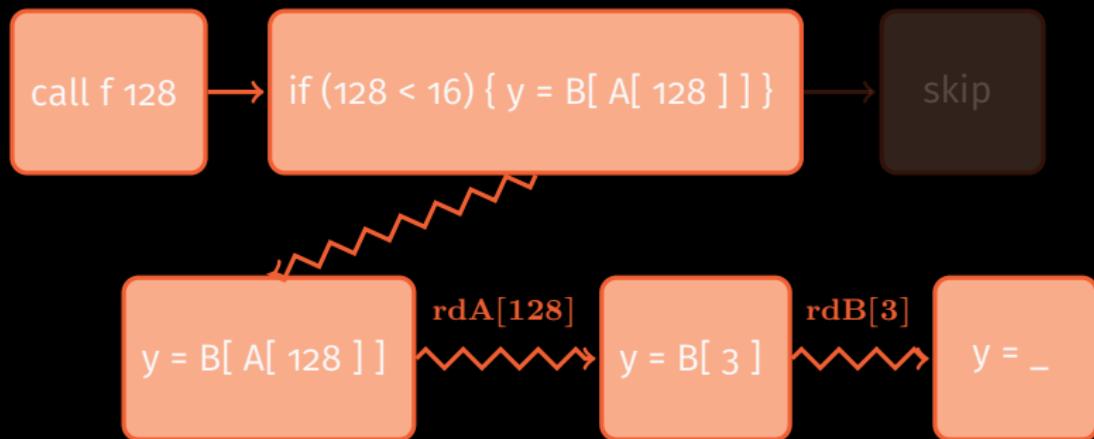
```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] }      // A.size=16, A[128]=3
```



Speculative Semantics 101

"Spectector..." S&P'20

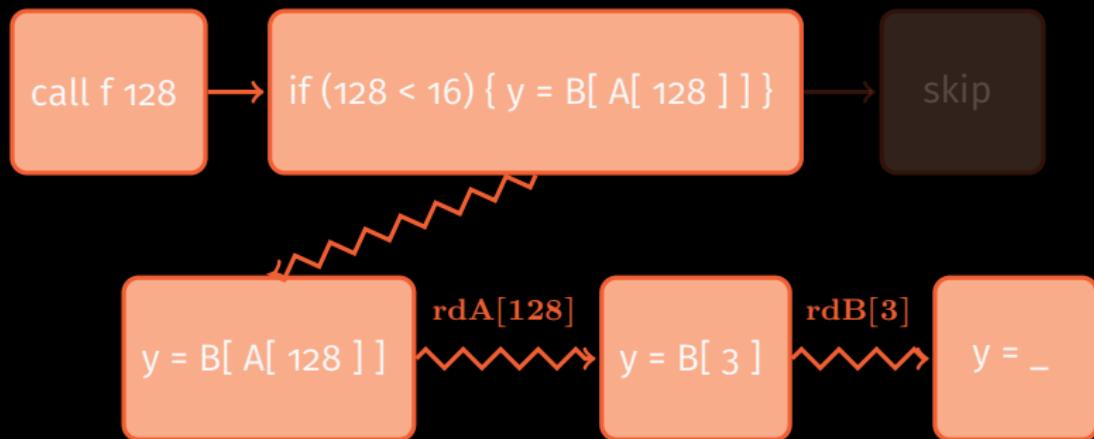
```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] }      // A.size=16, A[128]=3
```



Speculative Semantics 101

“Spectector ...” S&P'20

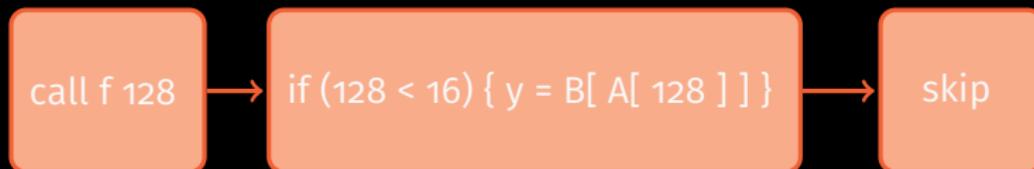
```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] }      // A.size=16, A[128]=3
```



Speculative Semantics 101

“Spectector ...” S&P'20

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] }      // A.size=16, A[128]=3
```



rdA[128]

rdB[3]

Speculative Safety (*SS*): Taint Tracking

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] } // A.size=16, A[128]=3
```

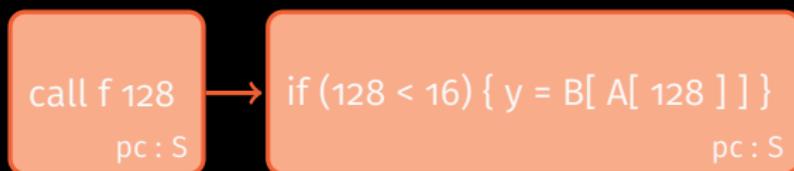
integrity lattice: $S \subset U$ $S \sqcap U = S$ U does not flow to S

```
call f 128  
pc : S
```

Speculative Safety (SS): Taint Tracking

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] } // A.size=16, A[128]=3
```

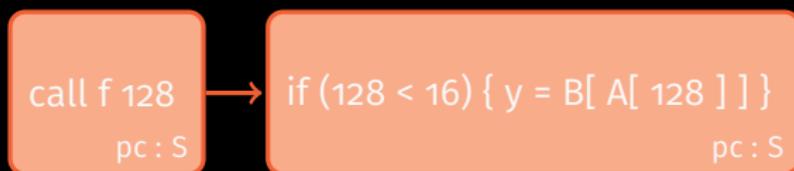
integrity lattice: $S \subset U$ $S \sqcap U = S$ U does not flow to S



Speculative Safety (SS): Taint Tracking

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] } // A.size=16, A[128]=3
```

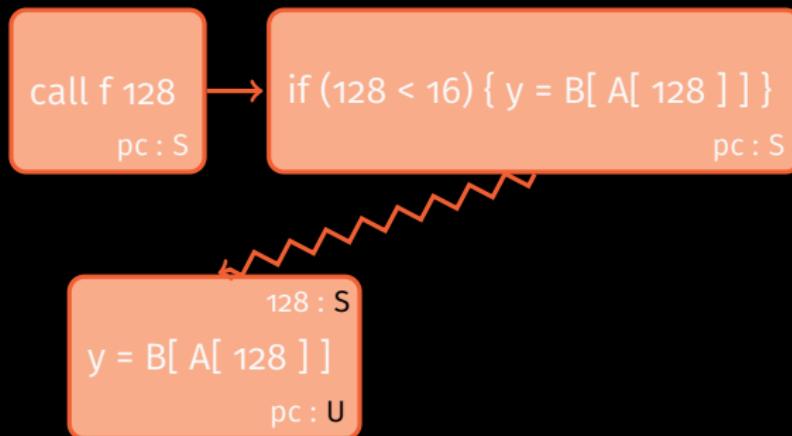
integrity lattice: $S \subset U$ $S \sqcap U = S$ U does not flow to S



Speculative Safety (SS): Taint Tracking

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] } // A.size=16, A[128]=3
```

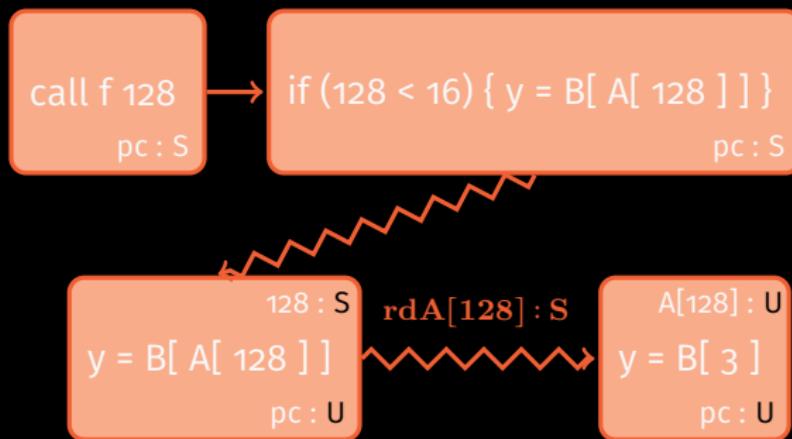
integrity lattice: $S \subset U$ $S \sqcap U = S$ U does not flow to S



Speculative Safety (*SS*): Taint Tracking

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] } // A.size=16, A[128]=3
```

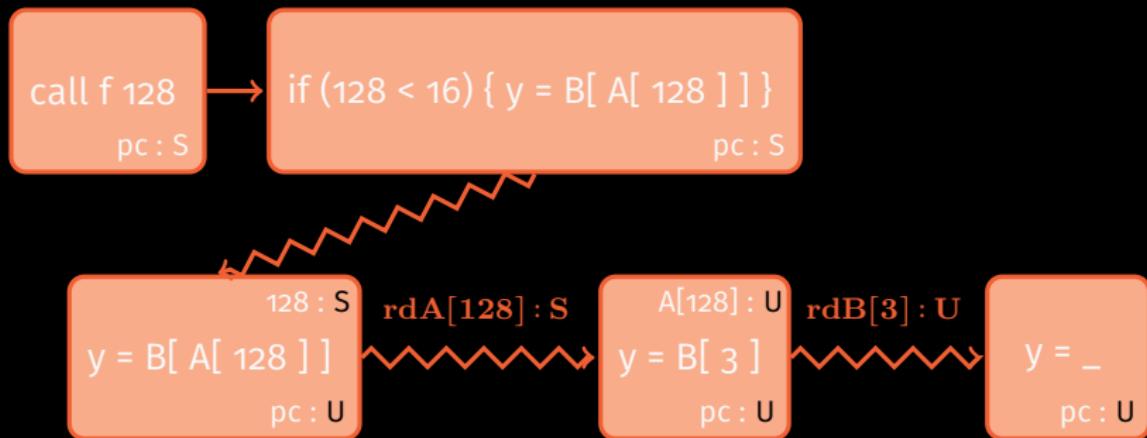
integrity lattice: $S \subset U$ $S \sqcap U = S$ U does not flow to S



Speculative Safety (*SS*): Taint Tracking

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] } // A.size=16, A[128]=3
```

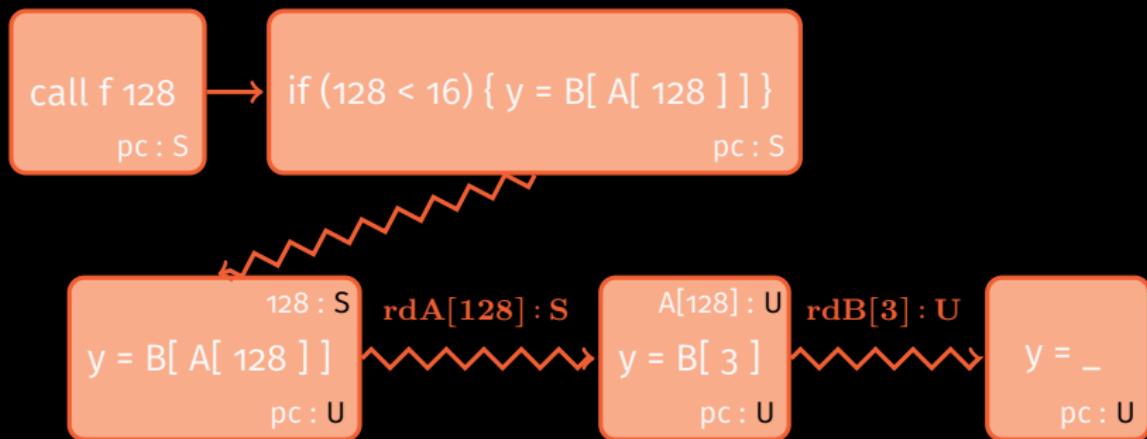
integrity lattice: $S \subset U$ $S \sqcap U = S$ U does not flow to S



Speculative Safety (*SS*): Taint Tracking

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] } // A.size=16, A[128]=3
```

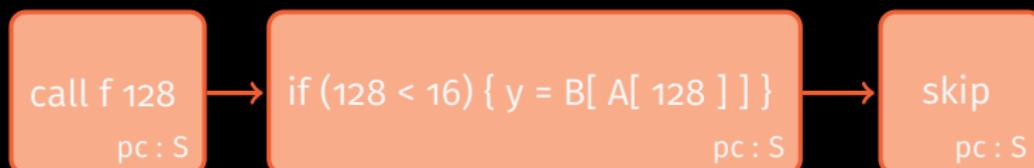
integrity lattice: $S \subset U$ $S \sqcap U = S$ U does not flow to S



Speculative Safety (SS): Taint Tracking

```
void f (int x) ↪ if (x < A.size) { y = B[ A[ x ] ] } // A.size=16, A[128]=3
```

integrity lattice: $S \subset U$ $S \sqcap U = S$ U does not flow to S



rdA[128] : S

rdB[3] : U

SS-Preserving Compiler: RSSC & RSSP

$\llbracket \cdot \rrbracket : \text{RSSP} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{blue}{A}. A[P] : SS \text{ then } \forall \textcolor{red}{A}. A[\llbracket P \rrbracket] : SS$

SS-Preserving Compiler: RSSC & RSSP

$\llbracket \cdot \rrbracket : \text{RSSP} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{blue}{A}.A[P] : SS \text{ then } \forall \textcolor{red}{A}.A[\llbracket P \rrbracket] : SS$

$\llbracket \cdot \rrbracket : \text{RSSC} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{red}{A}.A[\llbracket P \rrbracket] \rightsquigarrow m \text{ then } \exists \textcolor{blue}{A}.A[P] \rightsquigarrow m \approx m$

$\approx =$ same traces, plus S actions in m

SS-Preserving Compiler: RSSC & RSSP

$\llbracket \cdot \rrbracket : \text{RSSP} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{blue}{A}. A[P] : SS \text{ then } \forall \textcolor{red}{A}. A[\llbracket P \rrbracket] : SS$

$\llbracket \cdot \rrbracket : \text{RSSC} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{red}{A}. A[\llbracket P \rrbracket] \rightsquigarrow m \text{ then } \exists \textcolor{blue}{A}. A[P] \rightsquigarrow m \approx m$

$\approx =$ same traces, plus S actions in m

- RSSC & RSSP are equivalent

SS-Preserving Compiler: RSSC & RSSP

$\llbracket \cdot \rrbracket : \text{RSSP} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{blue}{A}. A[P] : SS \text{ then } \forall \textcolor{red}{A}. A[\llbracket P \rrbracket] : SS$

$\llbracket \cdot \rrbracket : \text{RSSC} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{red}{A}. A[\llbracket P \rrbracket] \rightsquigarrow m \text{ then } \exists \textcolor{blue}{A}. A[P] \rightsquigarrow m \approx m$

$\approx =$ same traces, plus S actions in m

- RSSC & RSSP are equivalent
- lfence : RSSC because it has no speculation ($pc : S$ always)

SS-Preserving Compiler: RSSC & RSSP

$\llbracket \cdot \rrbracket : \text{RSSP} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{blue}{A}.A[P] : SS \text{ then } \forall \textcolor{red}{A}.A[\llbracket P \rrbracket] : SS$

$\llbracket \cdot \rrbracket : \text{RSSC} \stackrel{\text{def}}{=} \text{if } \forall \textcolor{red}{A}.A[\llbracket P \rrbracket] \rightsquigarrow m \text{ then } \exists \textcolor{blue}{A}.A[P] \rightsquigarrow m \approx m$

$\approx =$ same traces, plus S actions in m

- RSSC & RSSP are equivalent
- lfence : RSSC because it has no speculation ($pc: S$ always)
- SLH : RSSC because masking taints as S

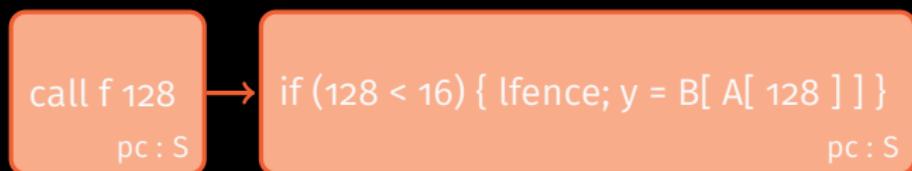
RSSC for lfence

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[.] = void f(int x) ↪ if(x < A.size){lfence; y = B[A[x]]}
```

```
call f 128  
pc : S
```

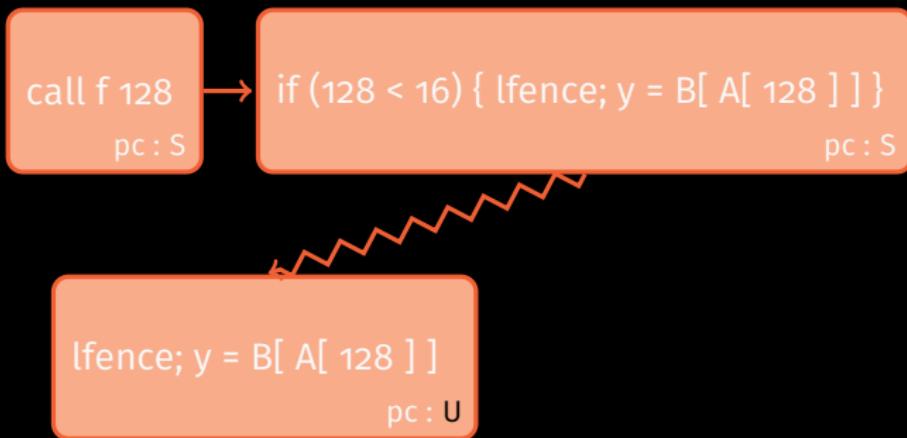
RSSC for lfence

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){lfence; y = B[A[x]]}
```



RSSC for lfence

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){lfence; y = B[A[x]]}
```



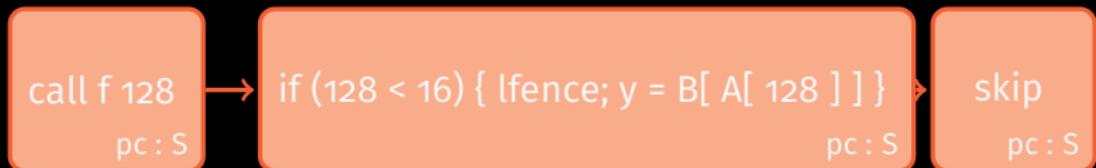
RSSC for lfence

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){lfence; y = B[A[x]]}
```



RSSC for lfence

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[.] = void f(int x) ↪ if(x < A.size){lfence; y = B[A[x]]}
```



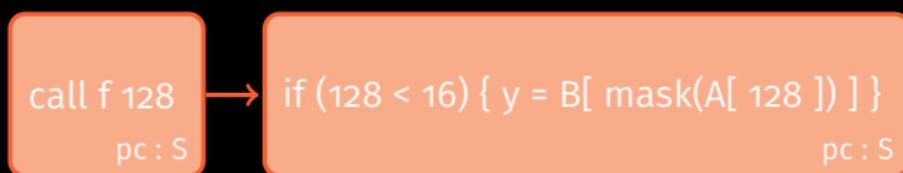
RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```

call f 128
pc : S

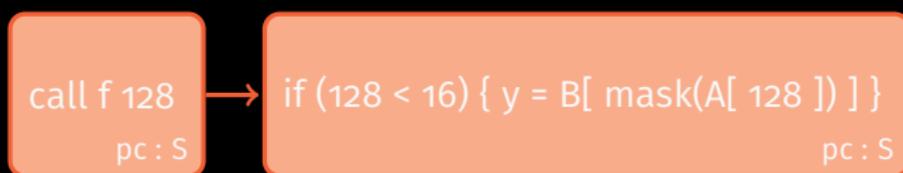
RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```



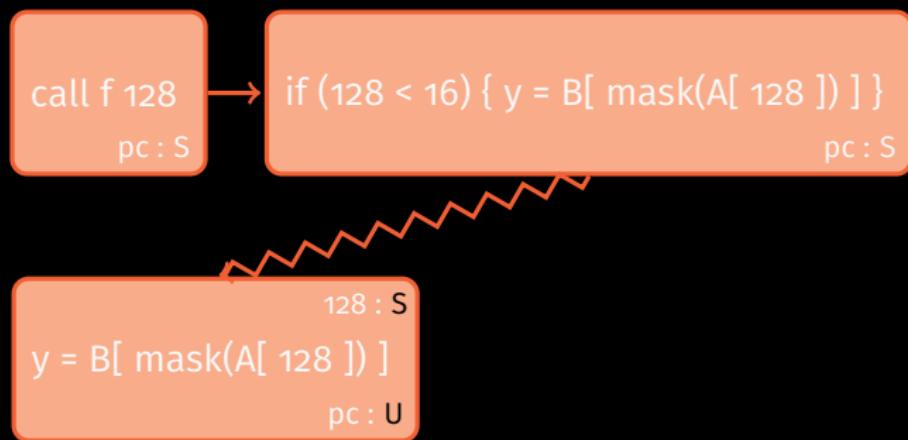
RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```



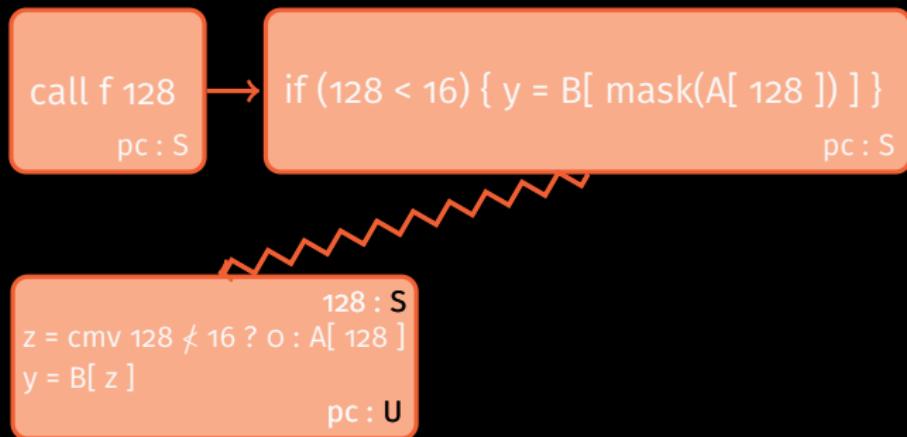
RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```



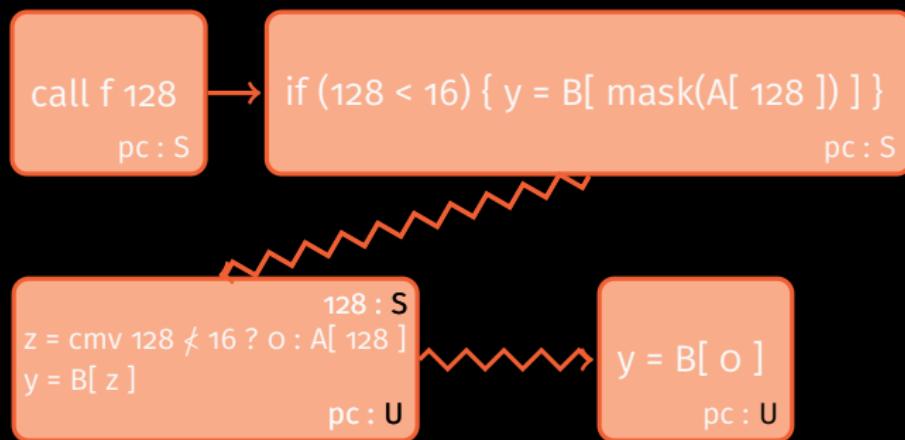
RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```



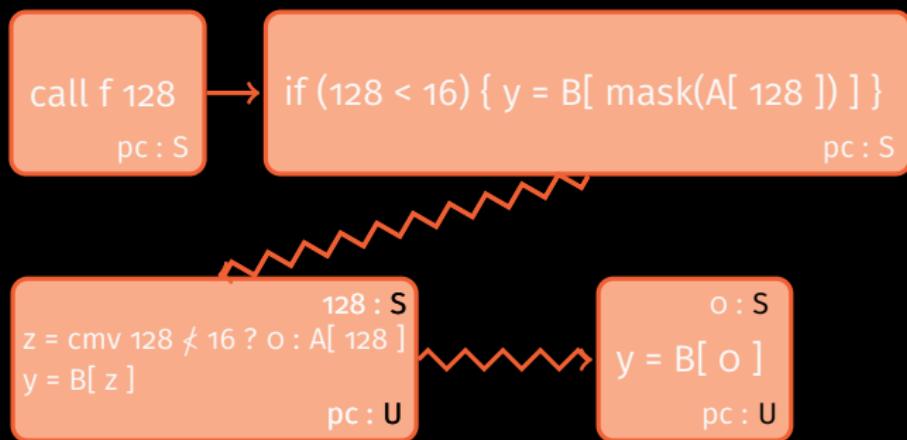
RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```



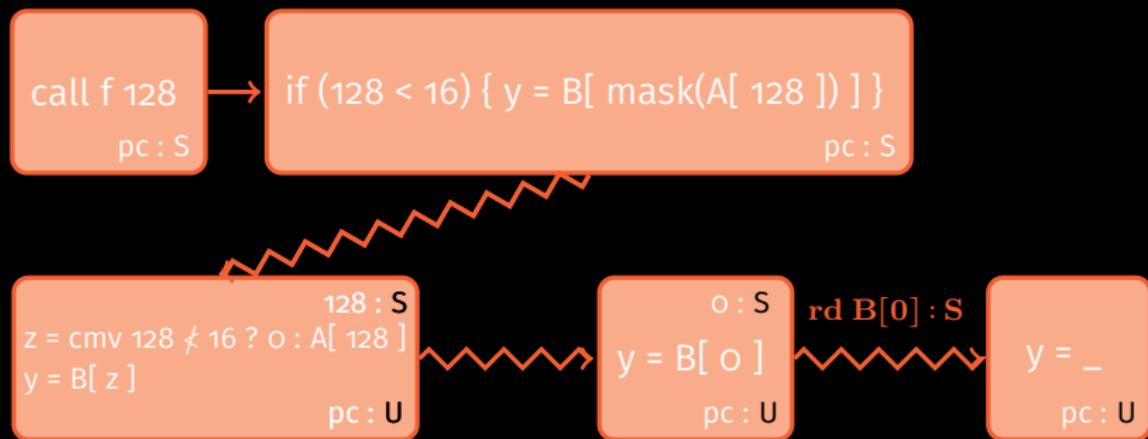
RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```



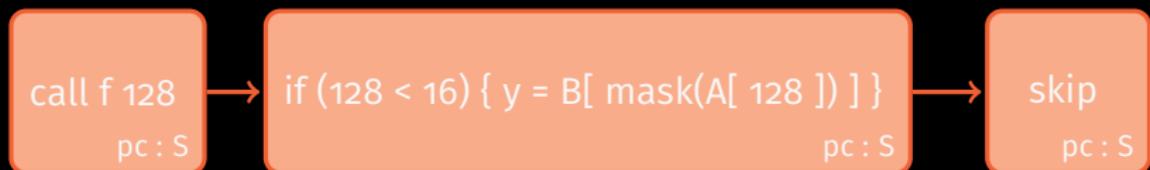
RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```



RSSC for SLH

```
void f(int x) ↪ if(x < A.size){y = B[A[x]]}      // A.size=16, A[128]=3  
[.] = void f(int x) ↪ if(x < A.size){y = B[mask(A[x])]}
```



rd B[0] : S

In the Paper

- theory scalable to Spectre v2, v4, v5

In the Paper

- theory scalable to Spectre v2, v4, v5
- 2 notions of *SS* for strong & weak absence of speculative leakage

In the Paper

- theory scalable to Spectre v2, v4, v5
- 2 notions of *SS* for strong & weak absence of speculative leakage
- proofs that RSSC means absence of leaks

In the Paper

- theory scalable to Spectre v2, v4, v5
- 2 notions of *SS* for strong & weak absence of speculative leakage
- proofs that RSSC means absence of leaks
- proof of security/insecurity for lfence, SLH (weakly), interprocedural SLH, SSLH

Future Outlook

What More?

What More?

- secure compilation for Spectre V2+
(w. Imdea, Cispa)

What More?

- secure compilation for Spectre V2+
(w. Imdea, Cispa)
- secure compilation to webassembly
(w. UCSD, Harvard, Cispa)

What More?

- secure compilation for Spectre V2+
(w. Imdea, Cispa)
- secure compilation to webassembly
(w. UCSD, Harvard, Cispa)
- secure compilation is universal
composability
(w. Stanford, Cispa)

What More?

- secure compilation for Spectre V2+
(w. Imdea, Cispa)
- secure compilation to webassembly
(w. UCSD, Harvard, Cispa)
- secure compilation is universal
composability
(w. Stanford, Cispa)
- secure compilation for linear languages
(w. Novi / FB)

What More?

- secure compilation for Spectre V2+
(w. Imdea, Cispa)
- secure compilation to webassembly
(w. UCSD, Harvard, Cispa)
- secure compilation is universal
composability
(w. Stanford, Cispa)
- secure compilation for linear languages
(w. Novi / FB)
- ...

Questions?

